


Meeting — Summary LLM


Bas Donkers, Dennis Fok, Finn-Ole Höner

May 13, 2024

- Gradient spikes first, then large step in decoder / encoder
- These steps seem to affect one hidden unit more strongly than the other
- Increase of Decoder Weight gradients seems to precede the increase of the -LL
- The other gradients seem to react to this jump
- Gradient-Decoder \rightarrow Step \rightarrow Other gradients \rightarrow Step ...



(a) Epoch 970-980



(a) Epoch 980-990

Figure 2: Spikes: Notice how the decoder gradient norm increases first, causing a large step and in-turn affecting the other gradients. The scales of the y-axis differ across all plots.

